IRON HACK

SUPERVISED LEARNING CLASSIFICATION

WHAT WILL WE COVER?

- Recap of what is supervised learning
- Core Concepts required
- Classification Algorithms
- Regression Algorithms

MACHINE LEARNING: SUPERVISED VS UNSUPERVISED

Whenever we have a defined output to predict it is a supervised ML problem

If we don't then we have a unsupervised learning problem

SUPERVISED: SHOULD I PLAY GOLF?

	OUTLOOK	TEMPERATURE	HUMIDITY	WINDY
0	Rainy	Hot	High	False
1	Rainy	Hot	High	True
2	Overcast	Hot	High	False
3	Sunny	Mild	High	False
4	Sunny	Cool	Normal	False
5	Sunny	Cool	Normal	True
6	Overcast	Cool	Normal	True
7	Rainy	Mild	High	False
8	Rainy	Cool	Normal	False
9	Sunny	Mild	Normal	False
10	Rainy	Mild	Normal	True
11	Overcast	Mild	High	True
12	Overcast	Hot	Normal	False
13	Sunny	Mild	High	True

The objective is to predict if based on the weather conditions of a particular day, we should go play Golf?

SUPERVISED: SHOULD I PLAY GOLF?

	OUTLOOK	TEMPERATURE	HUMIDITY	WINDY	PLAY GOLF
0	Rainy	Hot	High	False	No
1	Rainy	Hot	High	True	No
2	Overcast	Hot	High	False	Yes
3	Sunny	Mild	High	False	Yes
4	Sunny	Cool	Normal	False	Yes
5	Sunny	Cool	Normal	True	No
6	Overcast	Cool	Normal	True	Yes
7	Rainy	Mild	High	False	No
8	Rainy	Cool	Normal	False	Yes
9	Sunny	Mild	Normal	False	Yes
10	Rainy	Mild	Normal	True	Yes
11	Overcast	Mild	High	True	Yes
12	Overcast	Hot	Normal	False	Yes
13	Sunny	Mild	High	True	No

But how does my algorithm learn?

Based on past datapoints!!!

This is the "Supervised" component

These are called the labels

Supervised Learning is performed on labelled data

THE MACHINE LEARNING PROCESS

Labeled Data

Feature Engineering

Train-Test Split

Pick Algorithm & Hyperparameters

Optimize Parameters of Algorithm

Based on Training data

Evaluate Algorithm with appropriate metric

Based on Test data

Process of obtaining labeled data is tricky and many times the bottleneck of obtaining an efficient algorithm

THE MACHINE LEARNING PROCESS

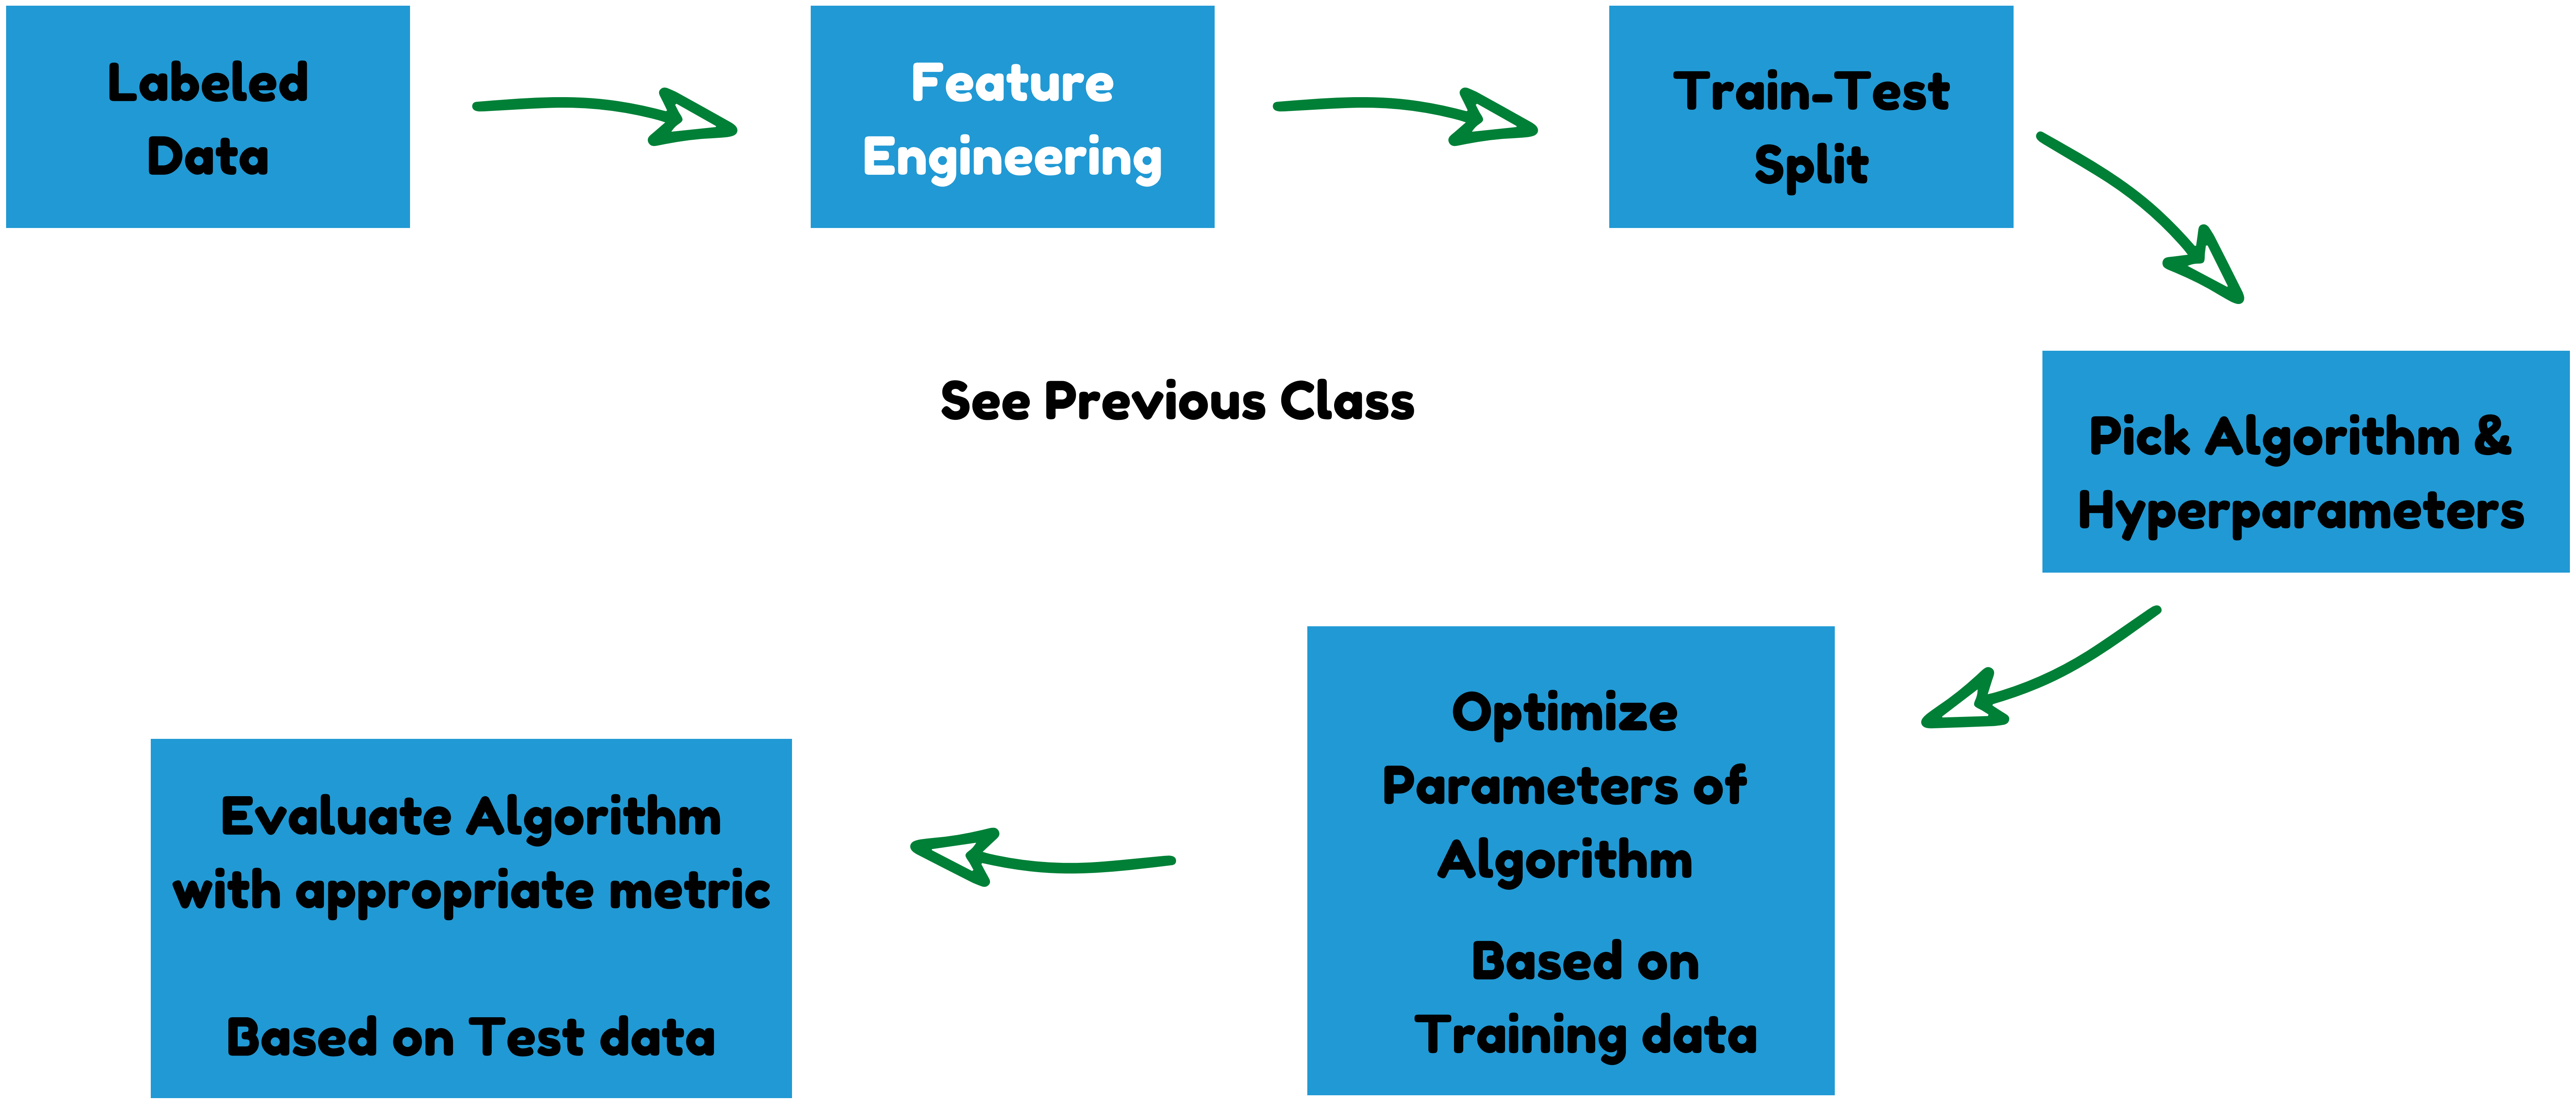

THE MACHINE LEARNING PROCESS

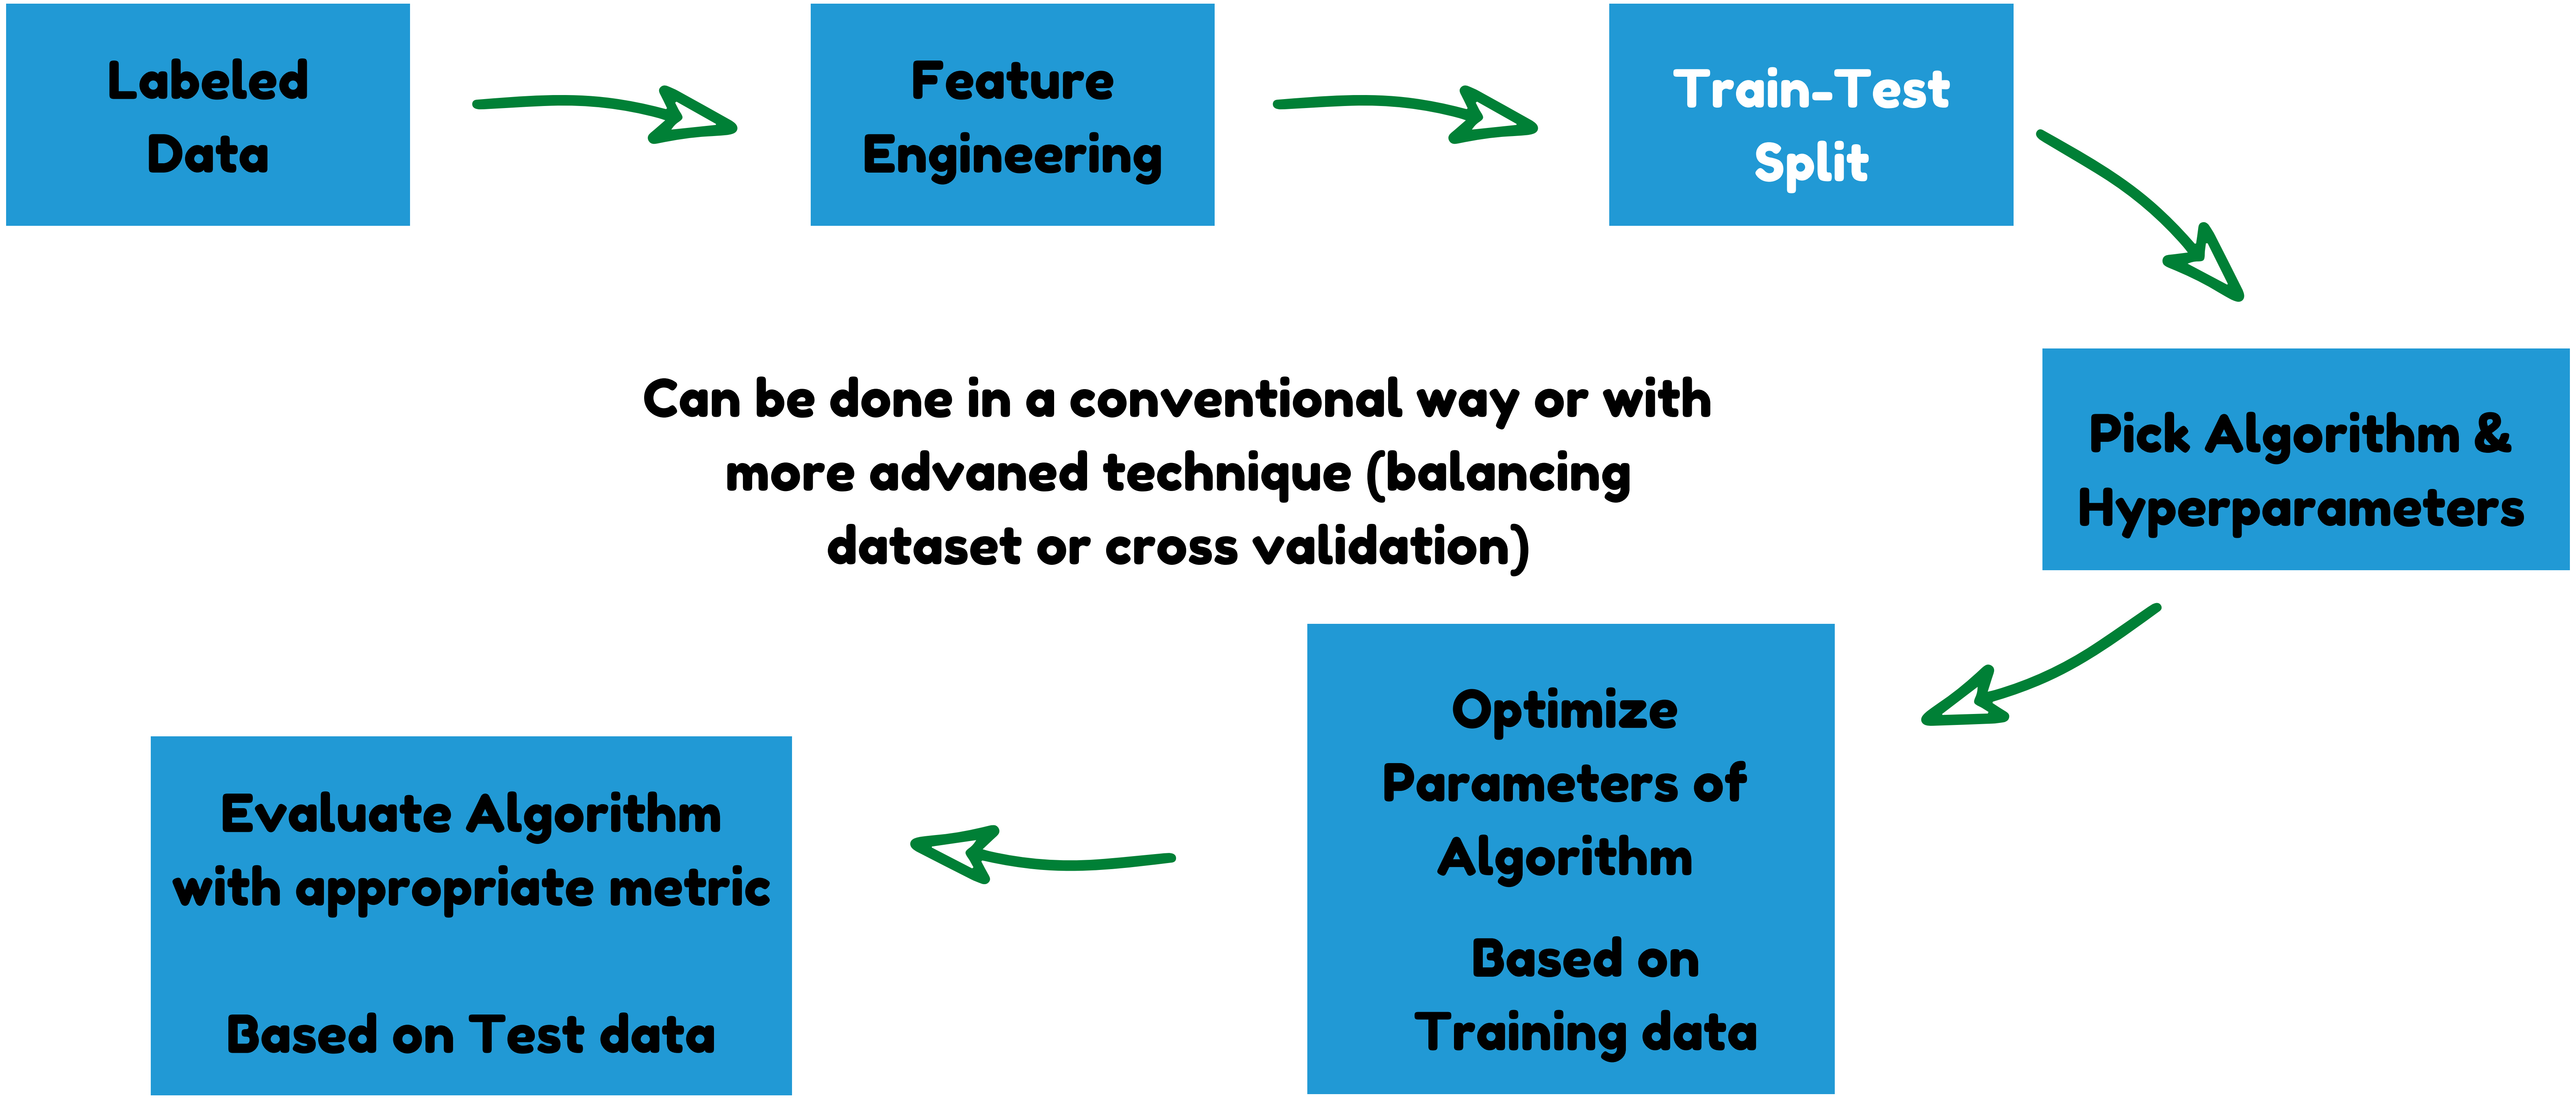

THE MACHINE LEARNING PROCESS

Labeled Data

Feature Engineering

Train-Test Split

Pick Algorithm & Hyperparameters

Optimize Parameters of Algorithm

Based on Training data

Evaluate Algorithm with appropriate metric

Based on Test data

With experience you will have intuition of which algorithm will be best suited for your situation. The hyperparameters have to be set before training

THE MACHINE LEARNING PROCESS

Labeled Data

Feature Engineering

Train-Test Split

Pick Algorithm & Hyperparameters

Optimize Parameters of Algorithm

Based on Training data

Evaluate Algorithm with appropriate metric

Based on Test data

Loss Function vs Cost Function
Attributes a Penalty to each incorrect prediction. Parameters are optimized to minimize this. This is also called "training"

THE MACHINE LEARNING PROCESS

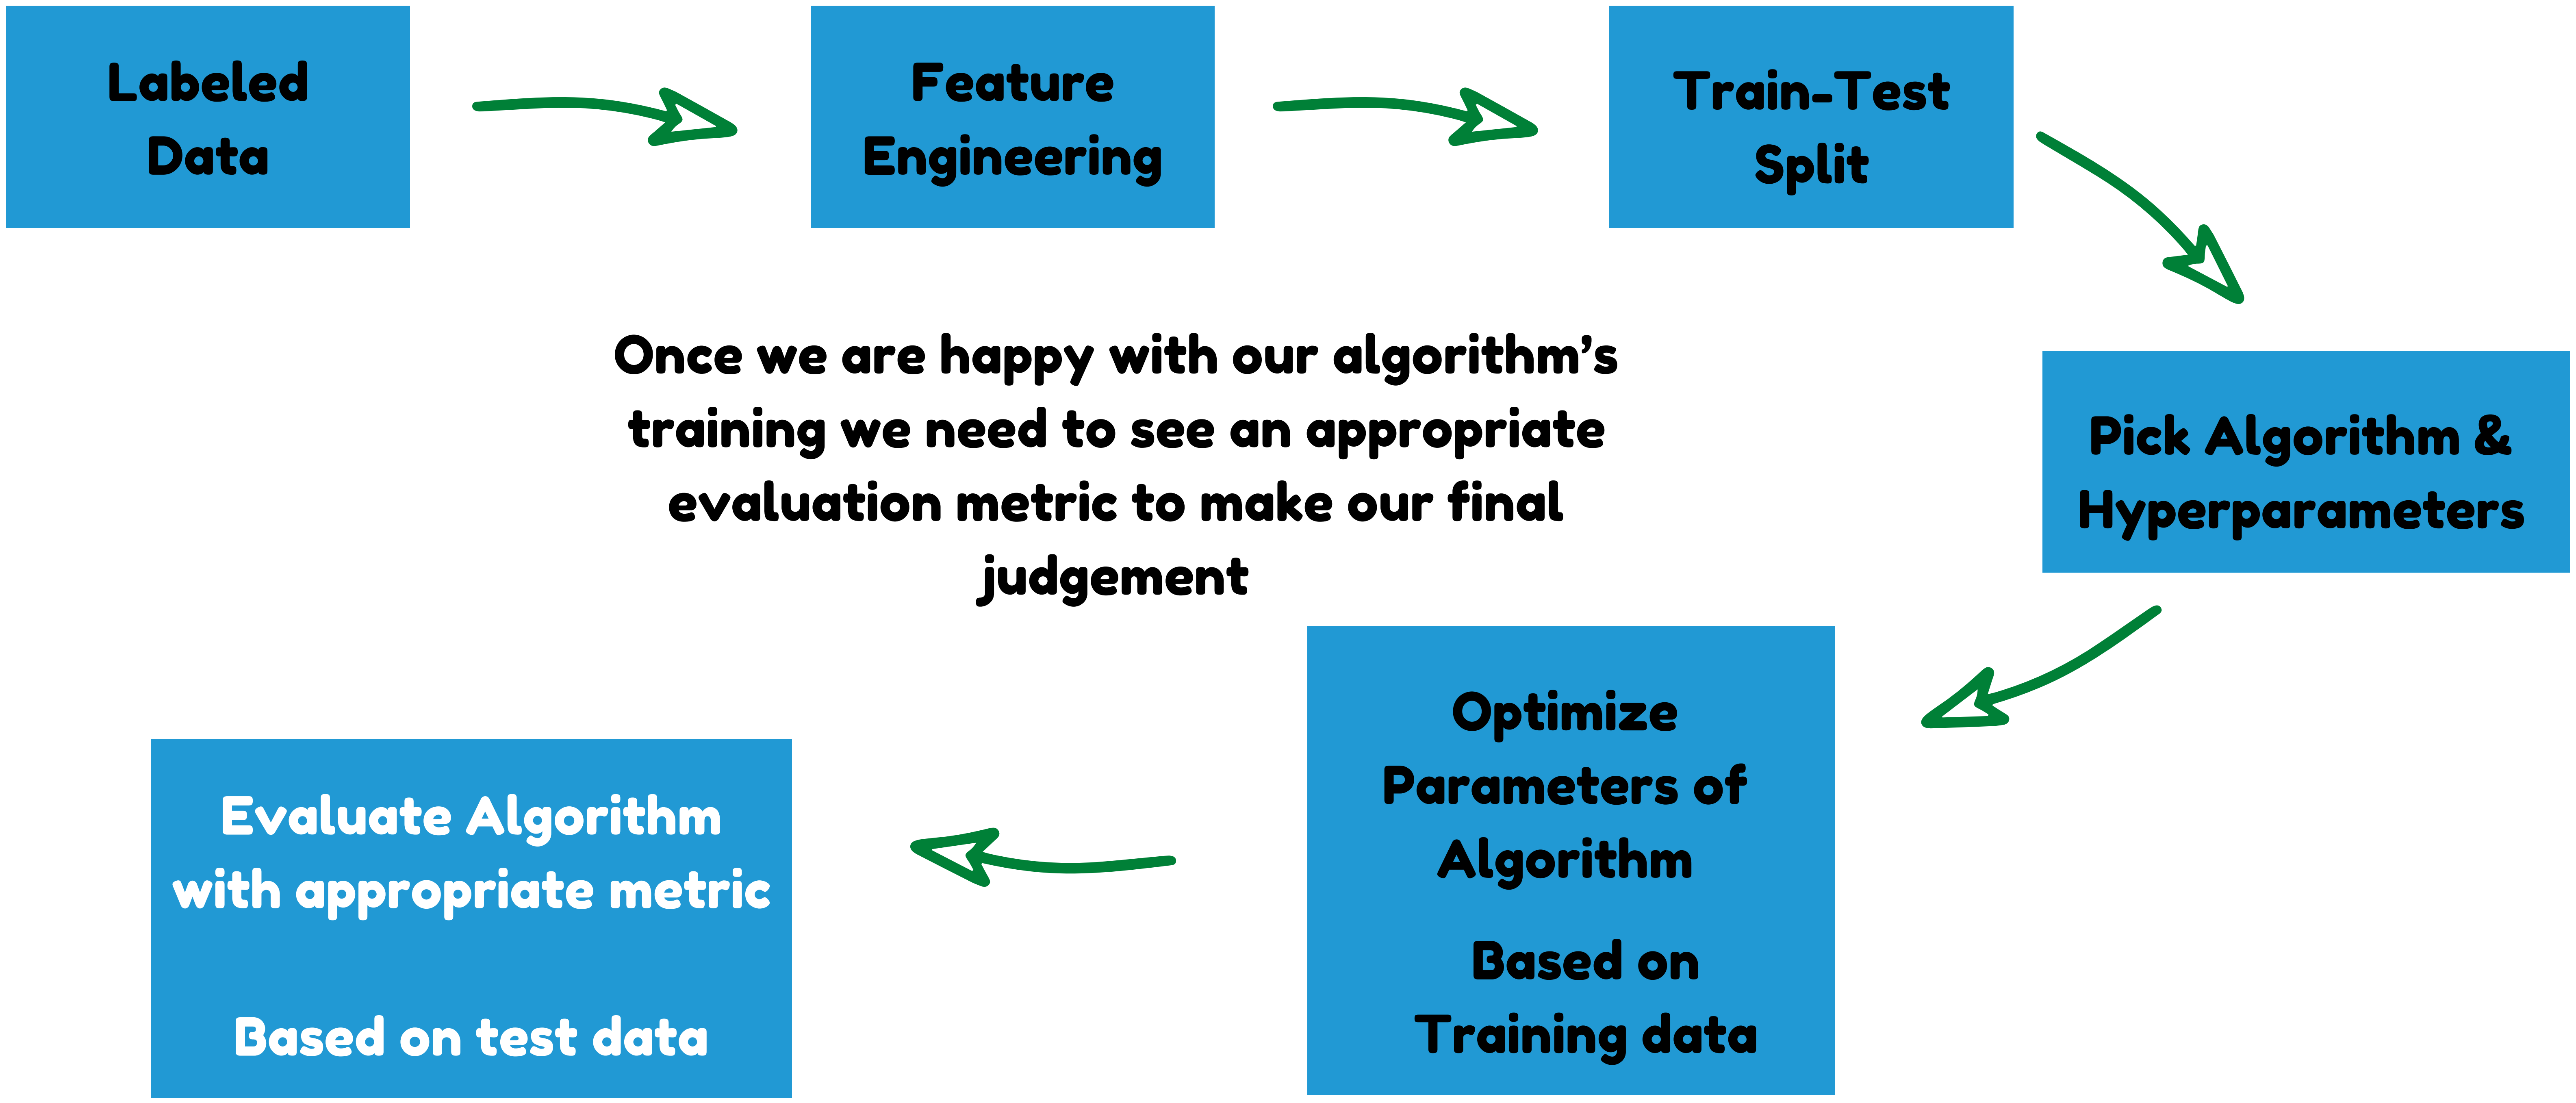

Once we are happy with our algorithm's training we need to see an appropriate evaluation metric to make our final judgement

THE MACHINE LEARNING PROCESS

Labeled Data

Feature Engineering

Train-Test Split

Pick Algorithm & Hyperparameters

Optimize Parameters of Algorithm

Based on Training data

Evaluate Algorithm with appropriate metric

Based on Test data

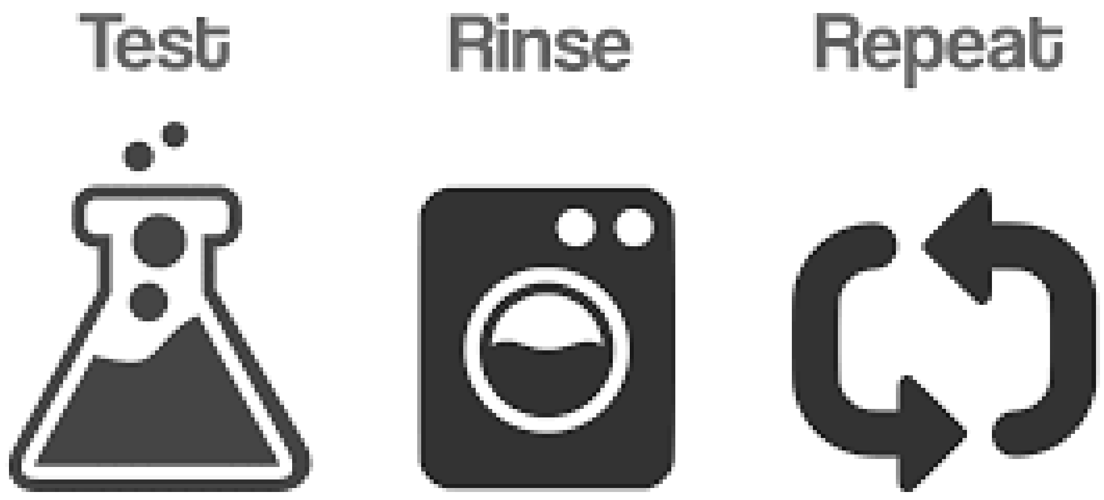

CLASSIFICATION ALGORITHMS

- KNN (Regression Friendly) (Multi Class)
- Logistic Regression
- Decision Trees (Regression Friendly) (Multi Class)
- Naive-Bayes Classifier (Multi Class)
- Support Vector Machine (SVM)

K NEAREST NEIGHBOURS

K- Nearest Neighbours (KNN)

The Copy-Cat of the Neighbours

Hyperparameters:

- K - # of neighbours
- Method of weight of neighbours

K NEAREST NEIGHBOURS

Example of Overfitting and training vs test tradeoff

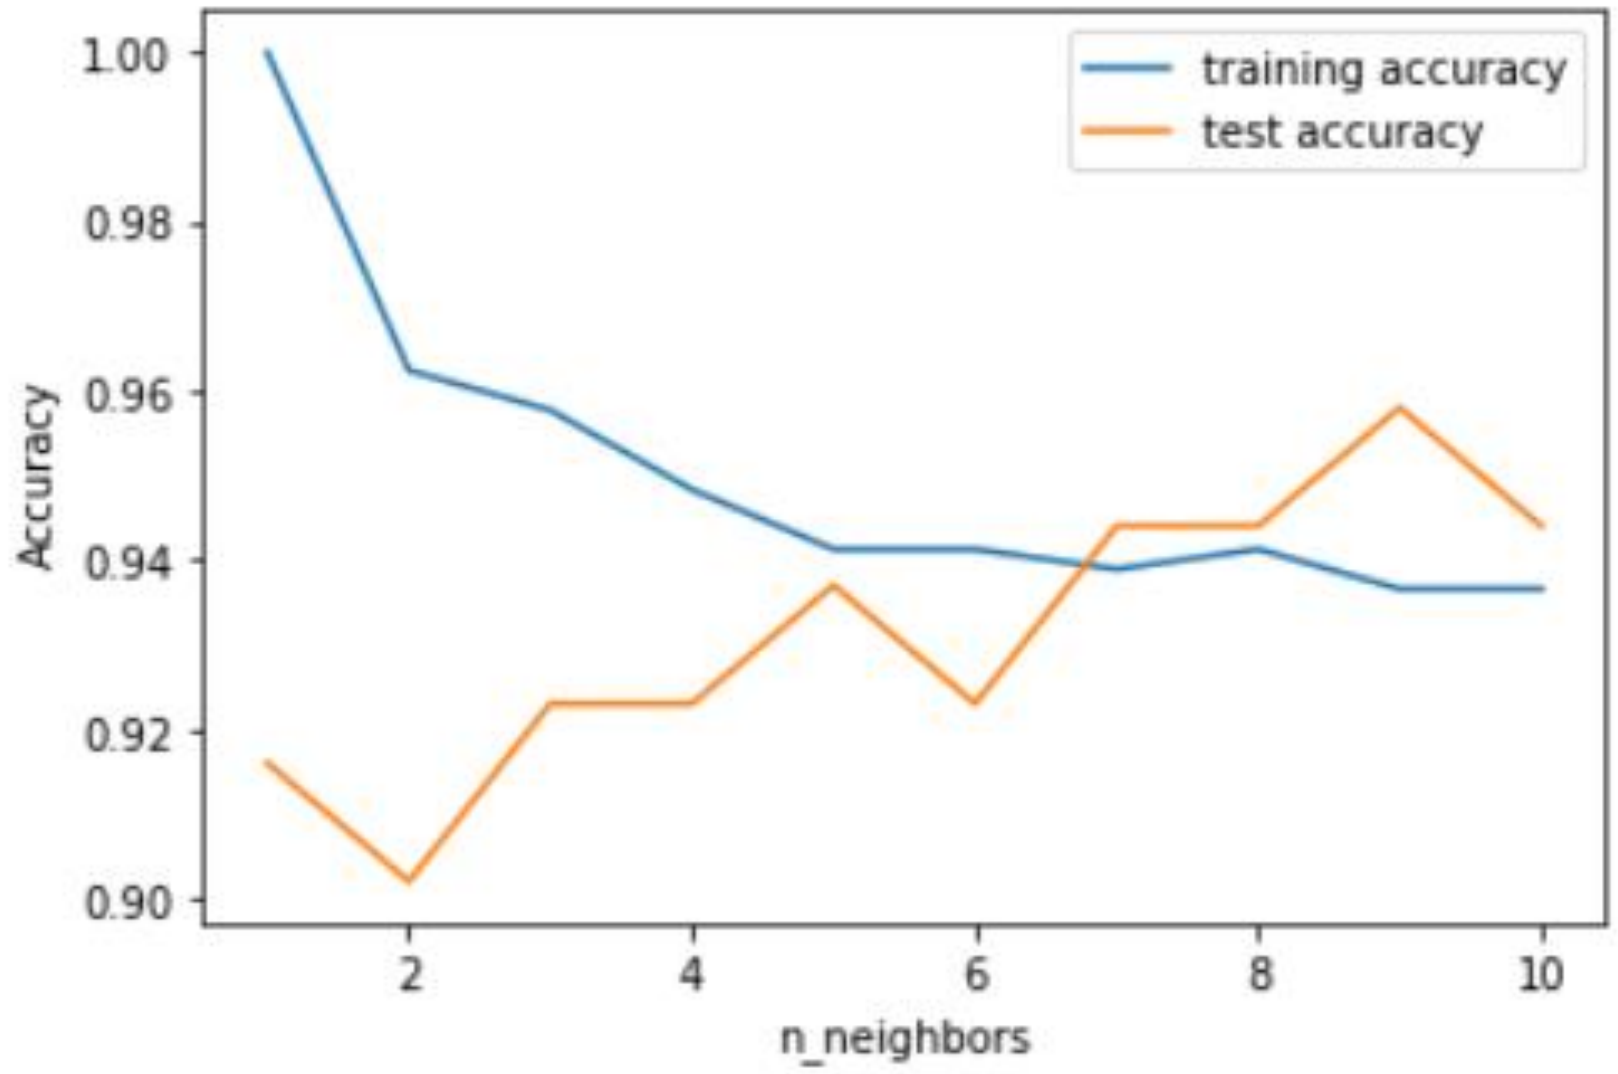

KNN - SUMMARY

+

No assumptions about data – useful, for example, for nonlinear data
Simple algorithm – to explain and understand/interpret

−

Requires all data to be stored in memory to compute
Can under perform with many variables
Very sensitive to scale

LOGISTIC REGRESSION

Binary Classification Algorithm

Takes in numerical inputs, takes a weighted sum of them and converts the result of that to a function that is "almost always" 0 or 1

Hyperparameter: Decision Boundary

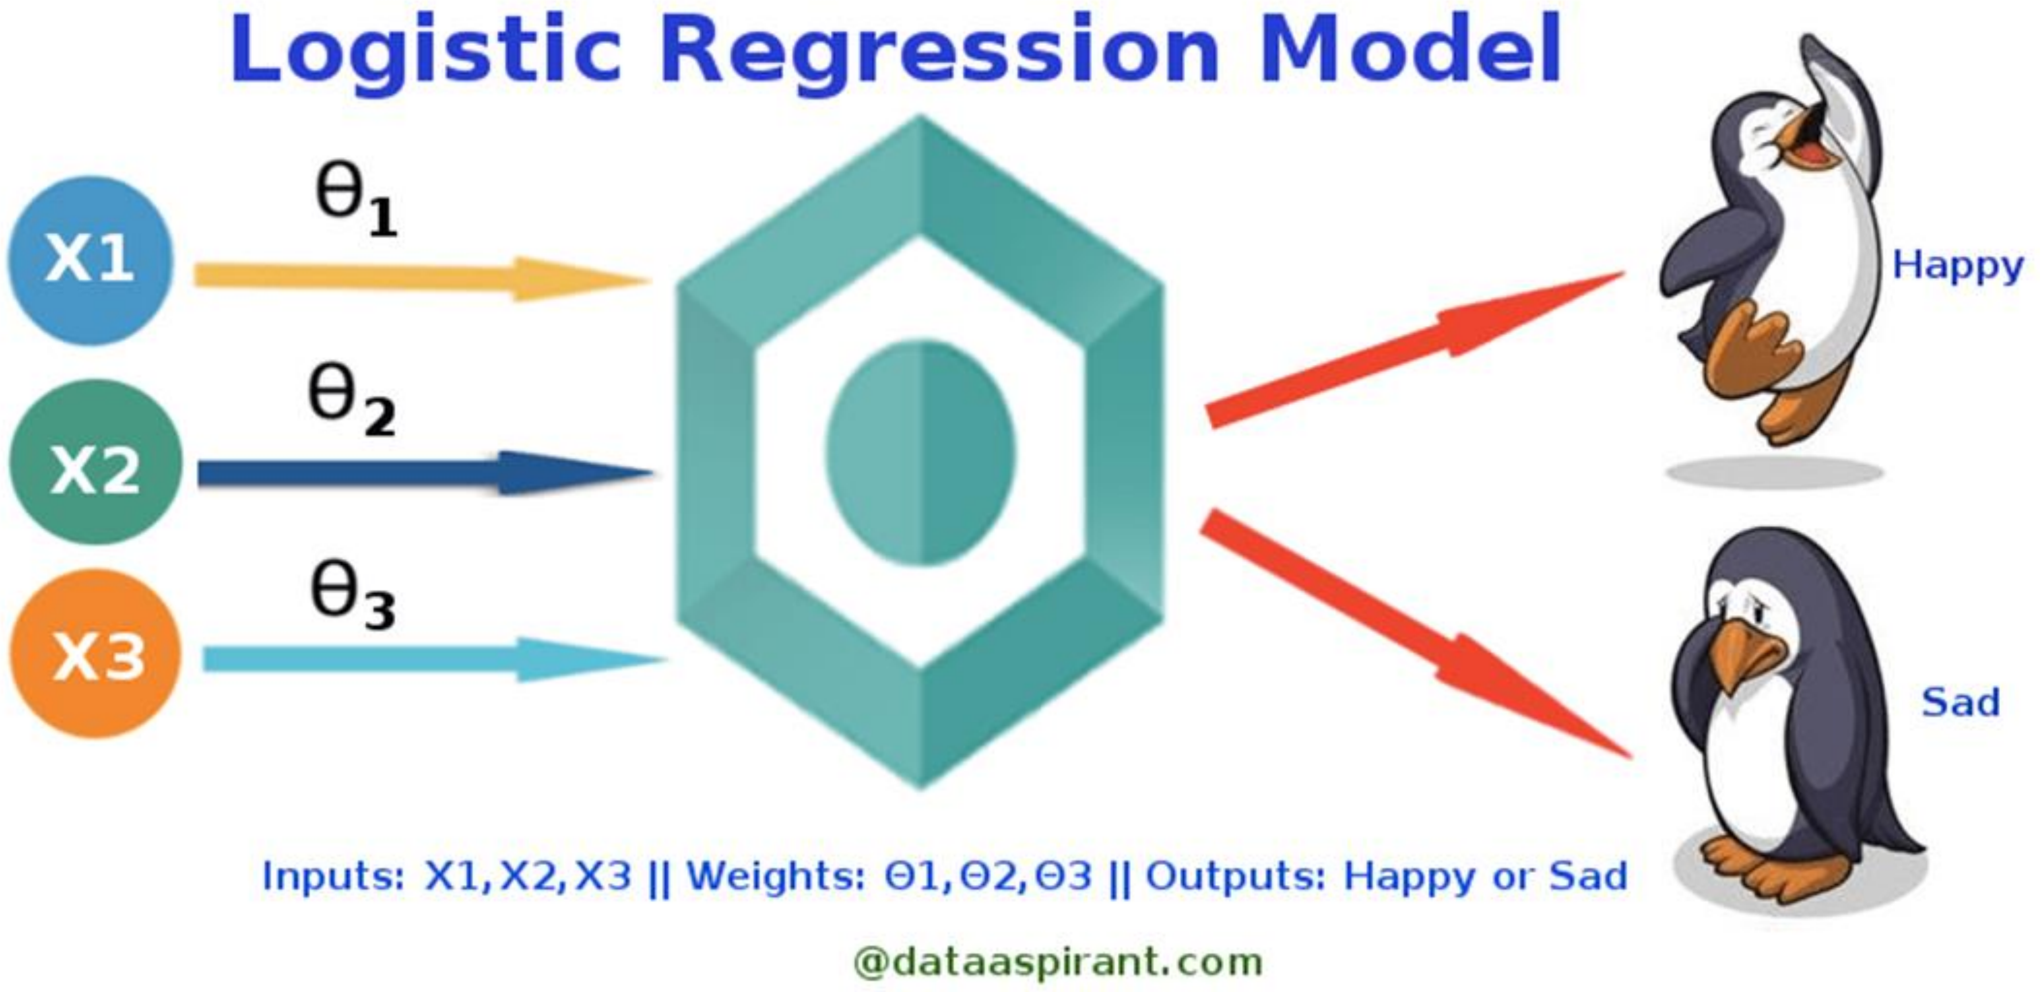

LOGISTIC REGRESSION

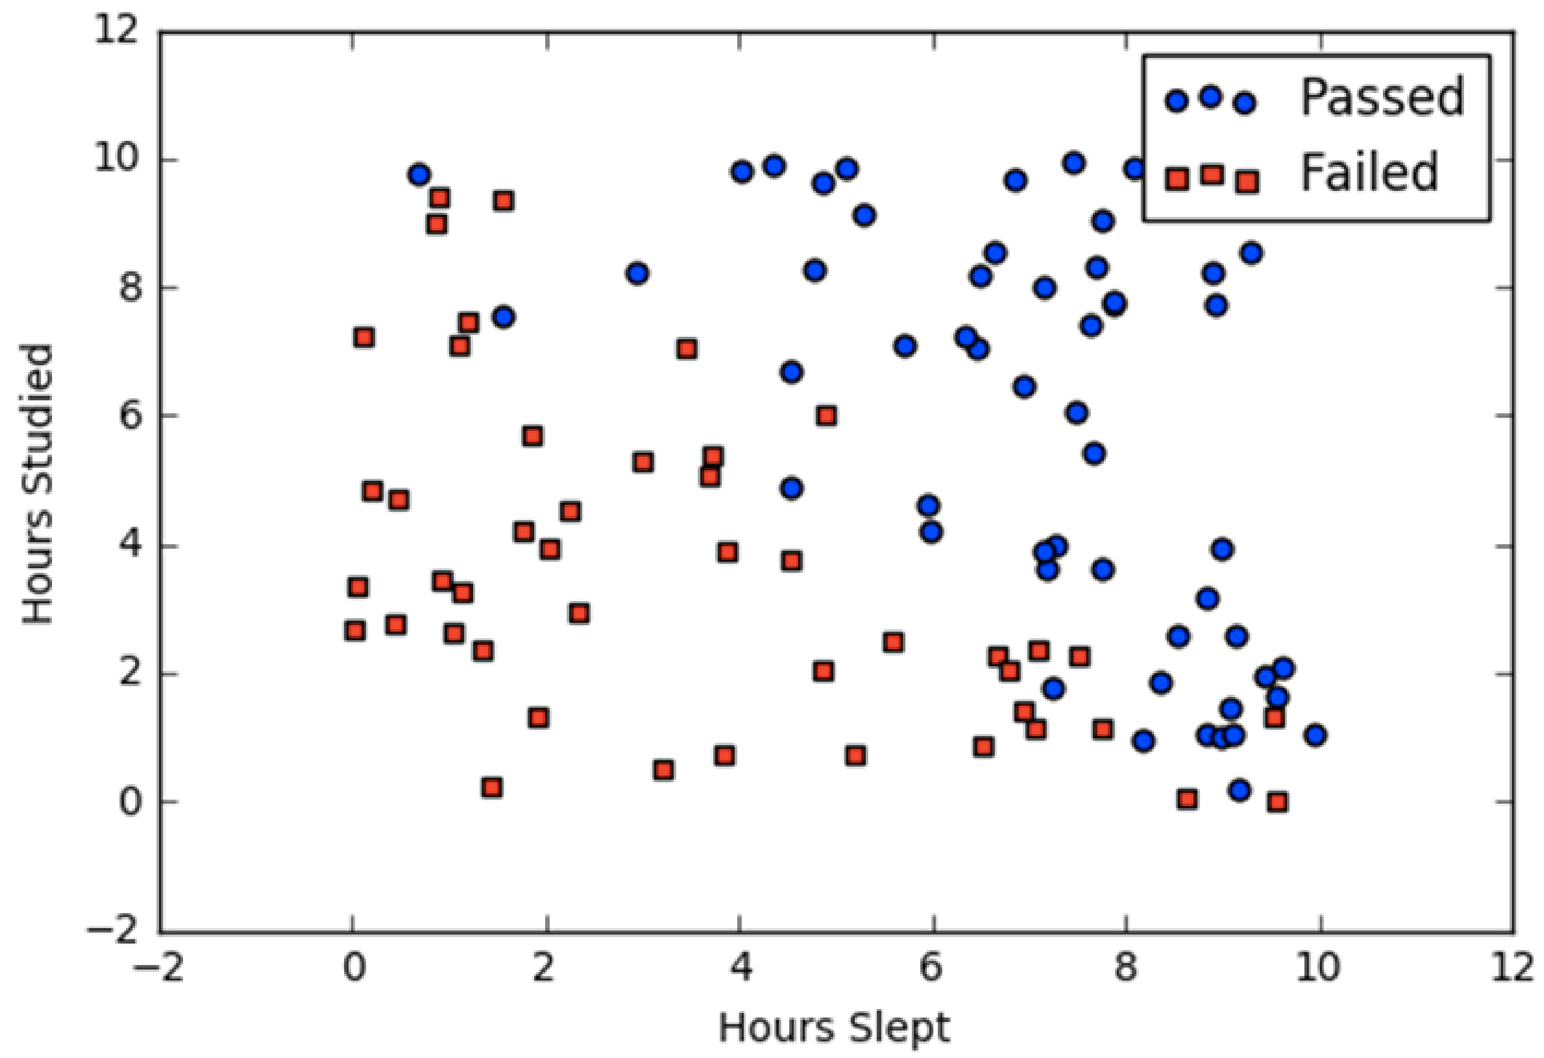

Studied	Slept	Passed
4.85	9.63	1
8.62	3.23	0
5.43	8.23	1
9.21	6.34	0

Steps:

- **Write usual linear combinations of input**
- **Pass result through sigmoid function**
- **Optimize parameters to minimize error**
- **Compare result to decision boundary**
- **Make Prediction**

LOGISTIC REGRESSION - LINEAR REGRESSION STEP

Studied	Slept	Passed
4.85	9.63	1
8.62	3.23	0
5.43	8.23	1
9.21	6.34	0

$$z = W_0 + W_1 Studied + W_2 Slept$$

We do not optimize yet the parameters of this linear regression and we don't yet compare z to anything

LOGISTIC REGRESSION - SIGMOID FUNCTION

This is where things become Nonlinear

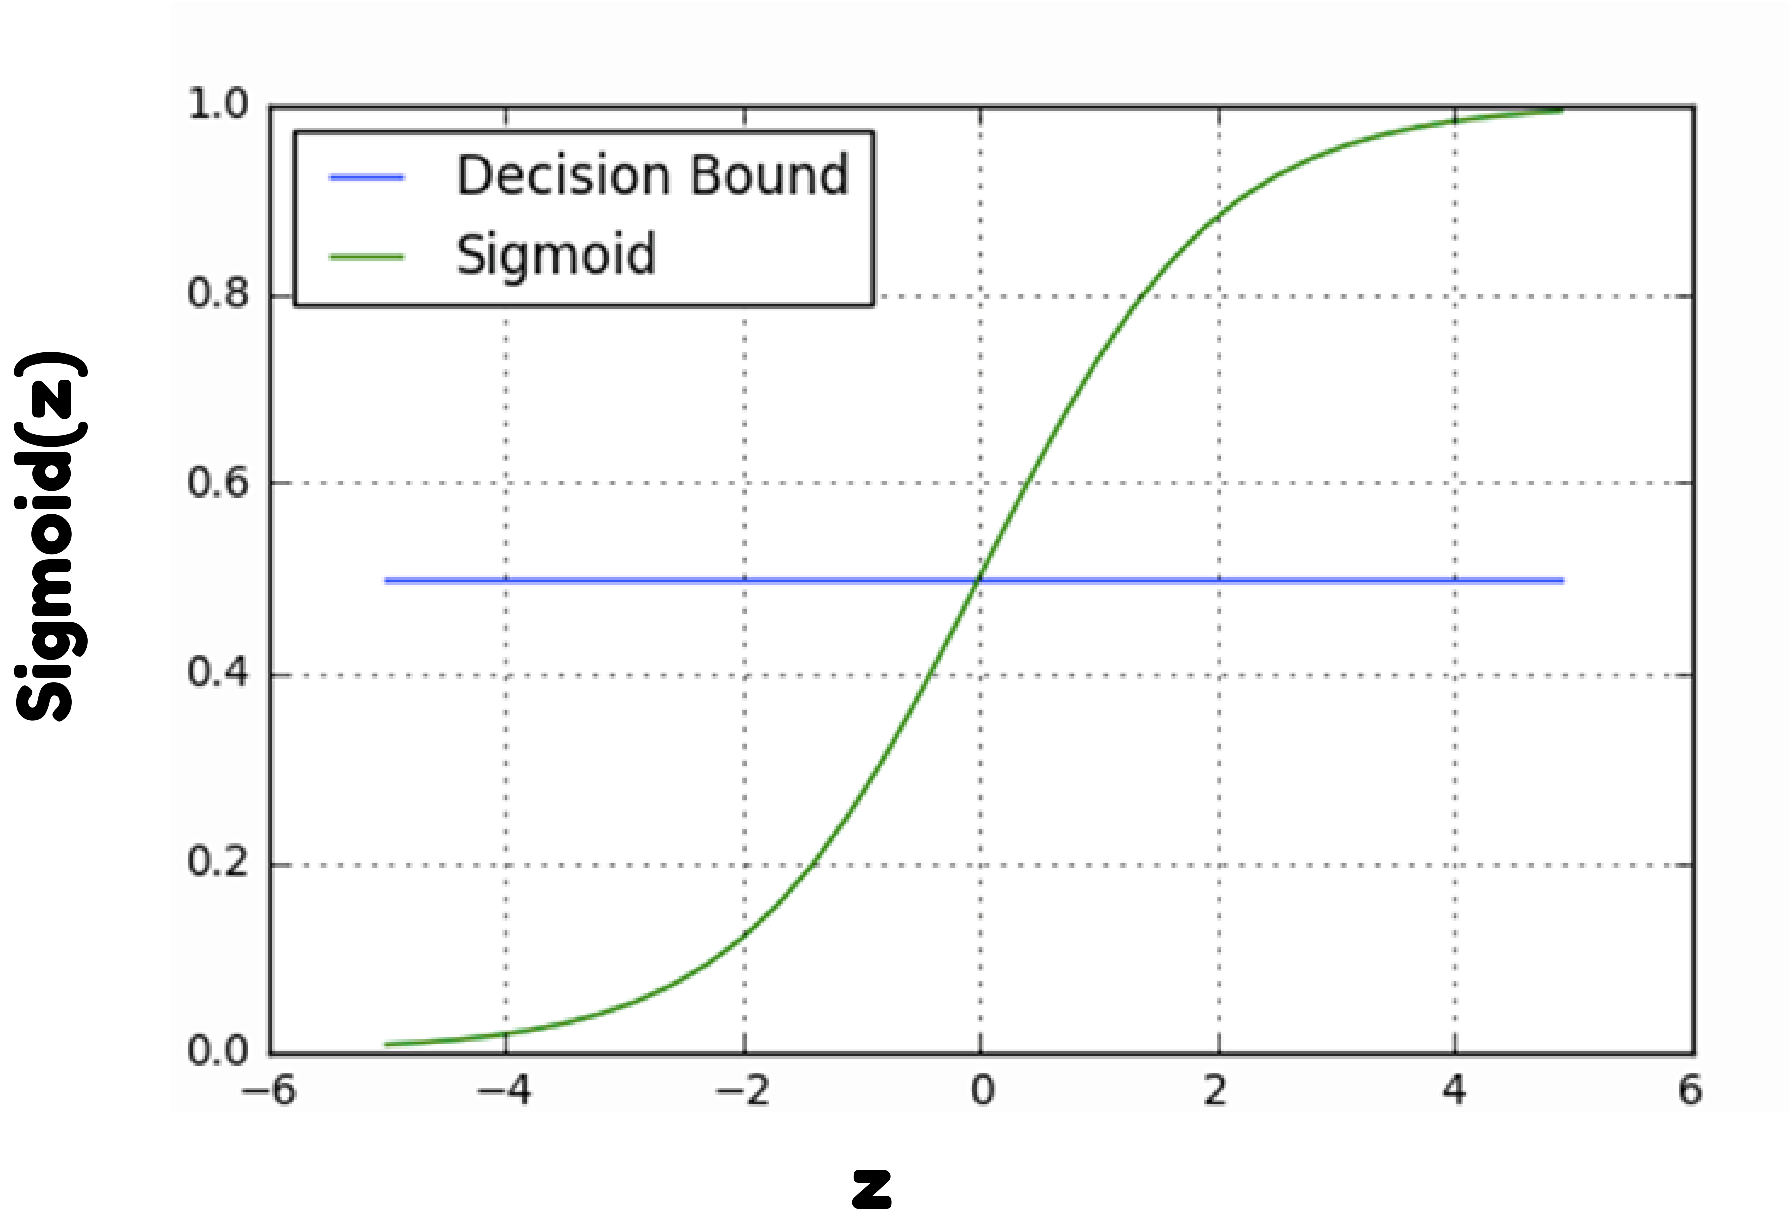

$$z = W_0 + W_1 Studied + W_2 Slept$$

$$f(z) = \frac{1}{1 + e^{-z}}$$

$$h(\vec{x}) = \frac{1}{1 + e^{-w_0 + w_1 x_1 + w_2 x_2}}$$

This function goes from 0 to 1 and typically does so relatively "sharply"
We <u>now</u> optimize W1 and W2 to "fit" the observations as well as possible

LOGISTIC REGRESSION - SUMMARY

+

Easy to implement
Trains quickly
Can be stacked to generate some really powerful models

−

Does not work as well for multi-class
Or at all for non-linearly separatable data

DECISION TREES

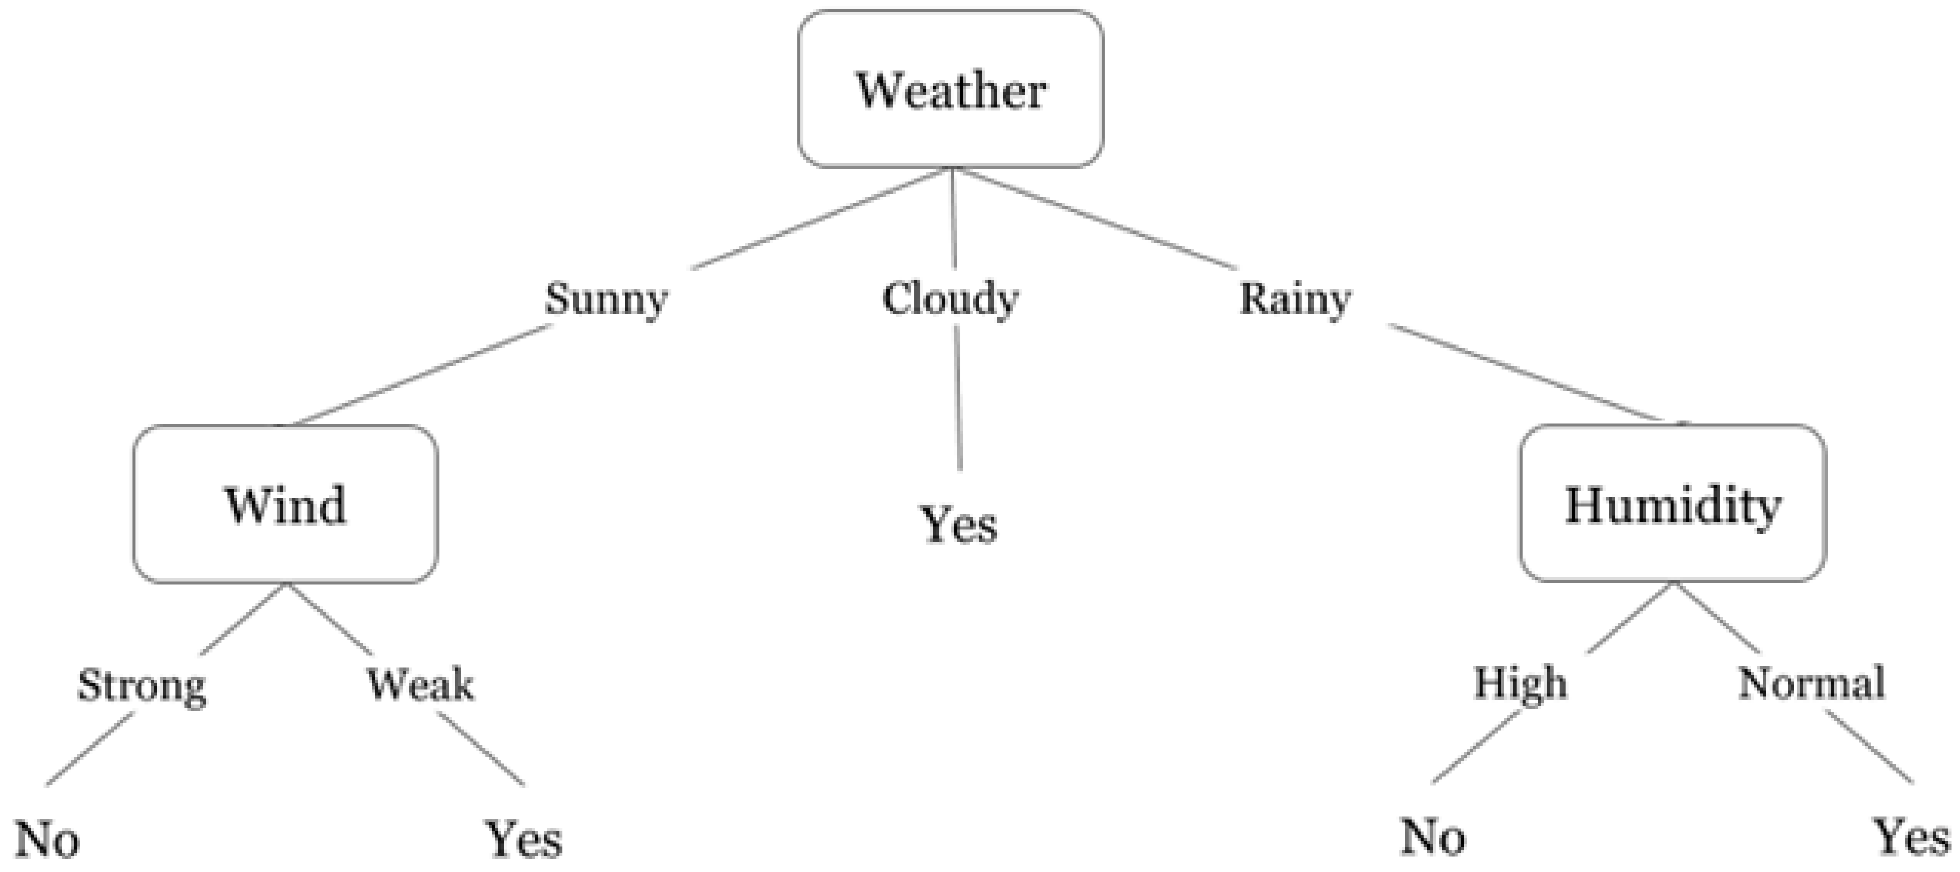

- Each level of the tree asks a "question" about a specific feature and divides the paths into several nodes

- number of levels is known as the tree depth

- The question you ask each level is the one that gives you the most accurate answer if you were to stop then

DECISION TREES

	OUTLOOK	TEMPERATURE	HUMIDITY	WINDY	PLAY GOLF
0	Rainy	Hot	High	False	No
1	Rainy	Hot	High	True	No
2	Overcast	Hot	High	False	Yes
3	Sunny	Mild	High	False	Yes
4	Sunny	Cool	Normal	False	Yes
5	Sunny	Cool	Normal	True	No
6	Overcast	Cool	Normal	True	Yes
7	Rainy	Mild	High	False	No
8	Rainy	Cool	Normal	False	Yes
9	Sunny	Mild	Normal	False	Yes
10	Rainy	Mild	Normal	True	Yes
11	Overcast	Mild	High	True	Yes
12	Overcast	Hot	Normal	False	Yes
13	Sunny	Mild	High	True	No

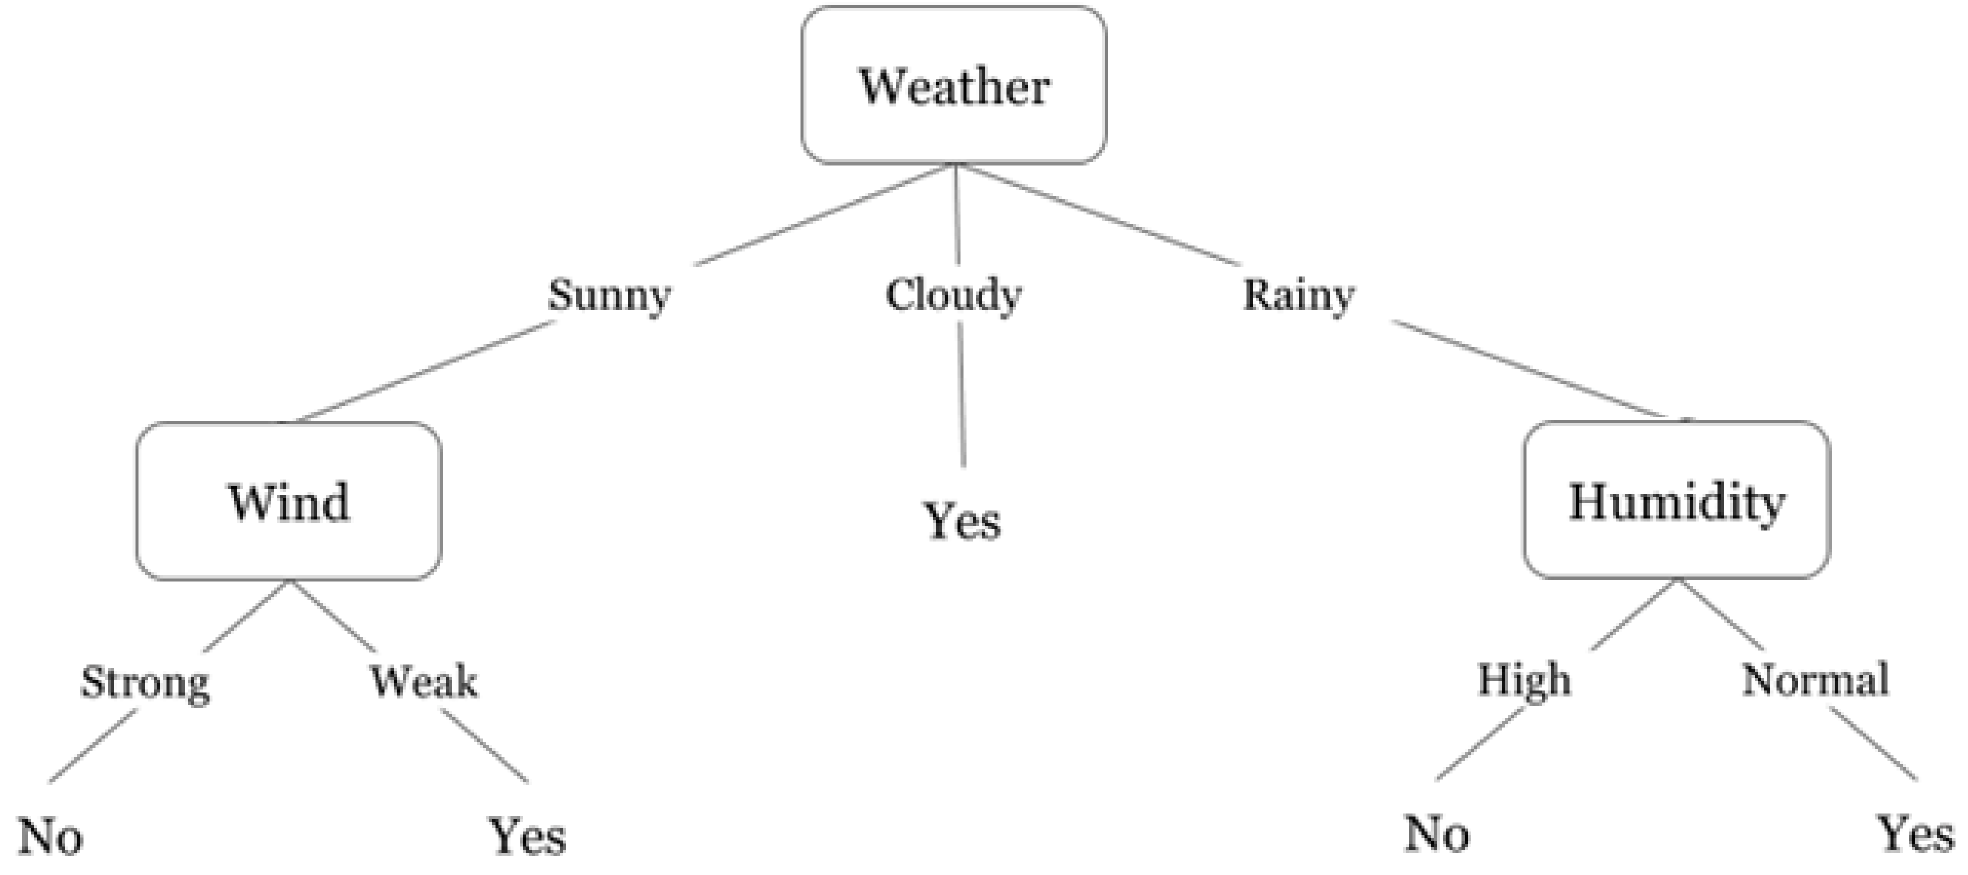

DECISION TREES - SUMMARY

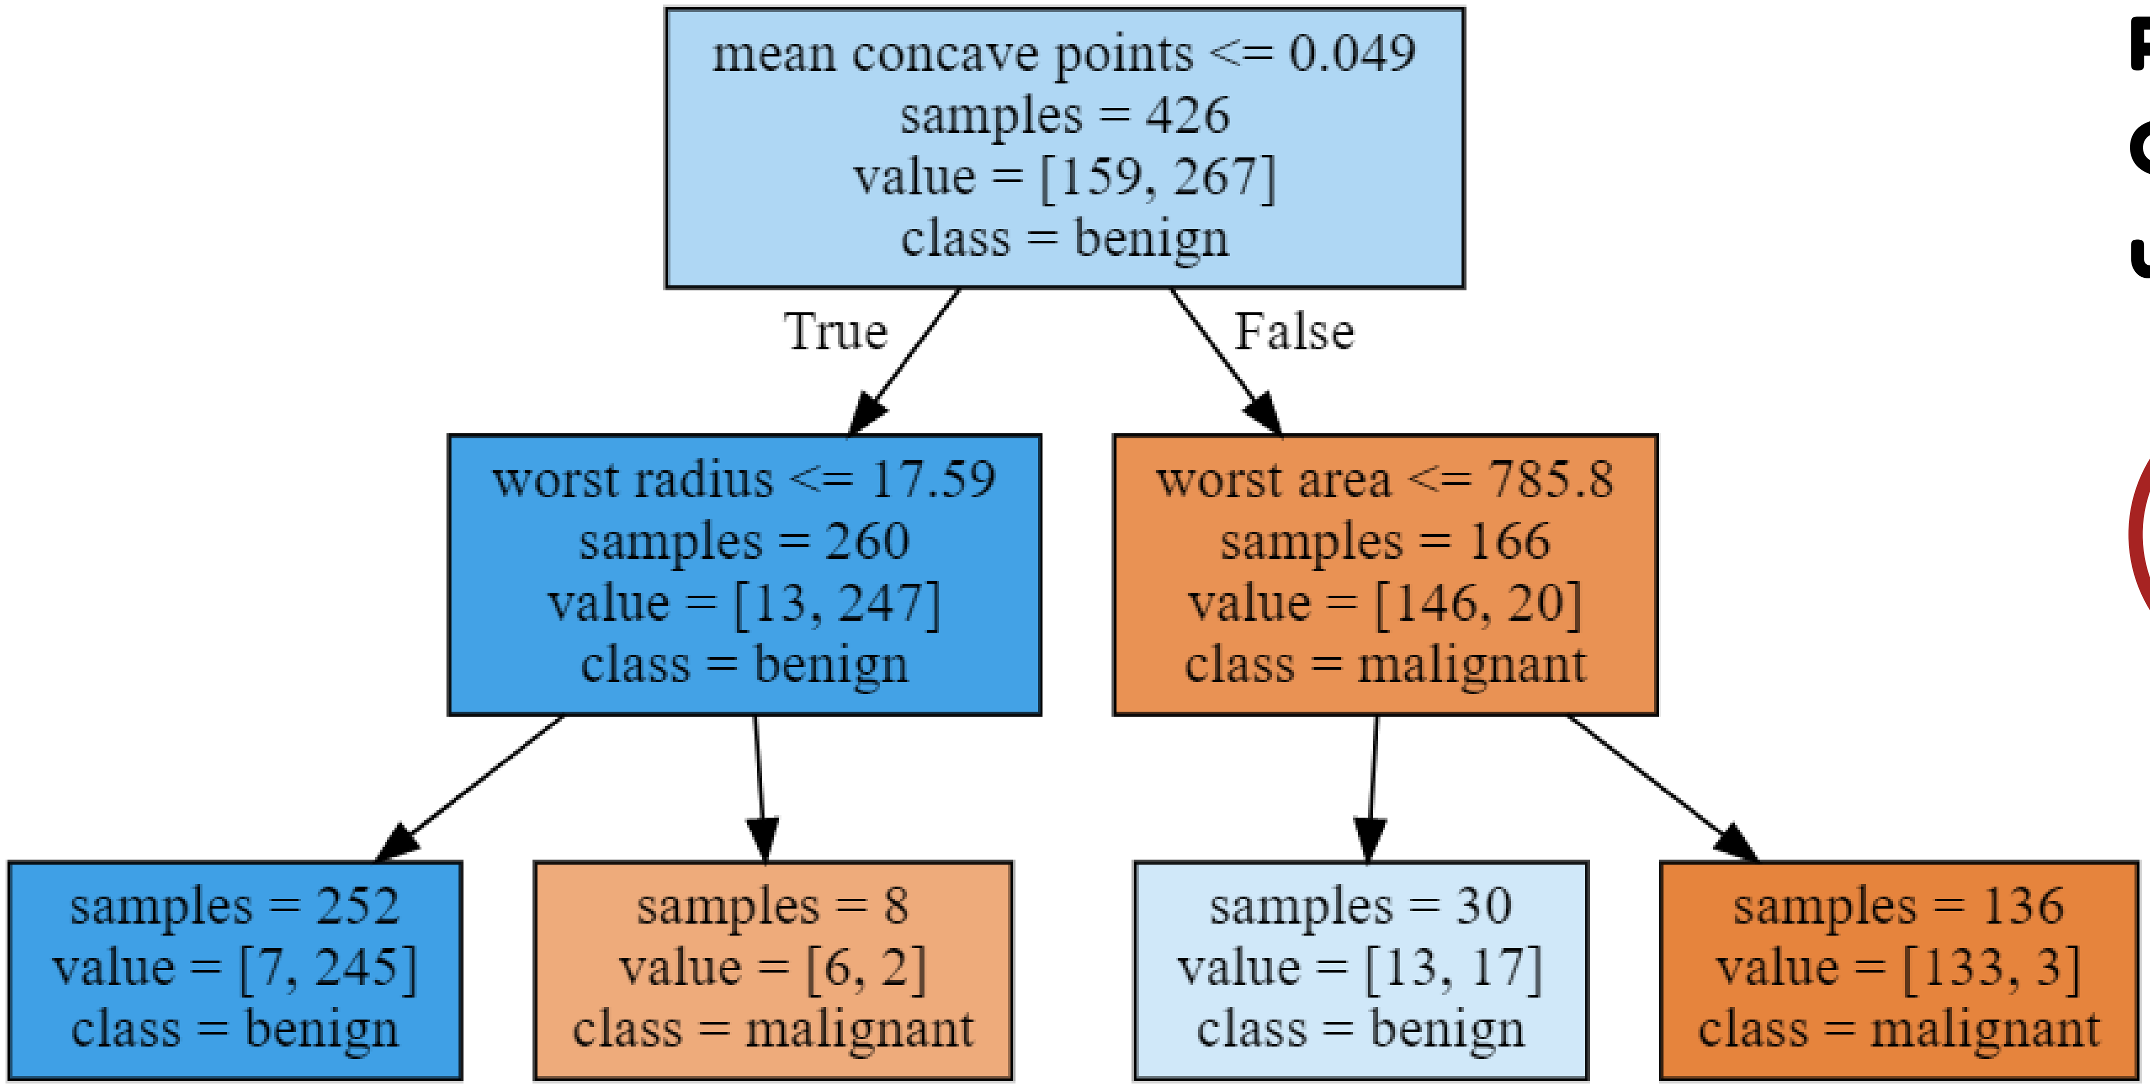

+

Very robust to different types of features, including NaNs, categorical, numerical

Handles non-linearity

Robust to outliers

Can be stacked into arguably the most useful models in modern ML

−

Very quickly overfit if hyperparameters are not controlled

Generalization Problems: what happens if feature outside possible ranges appears?

Online Learning problem: quickly change when re-trained

SUPPORT VECTOR MACHINE (SVM)

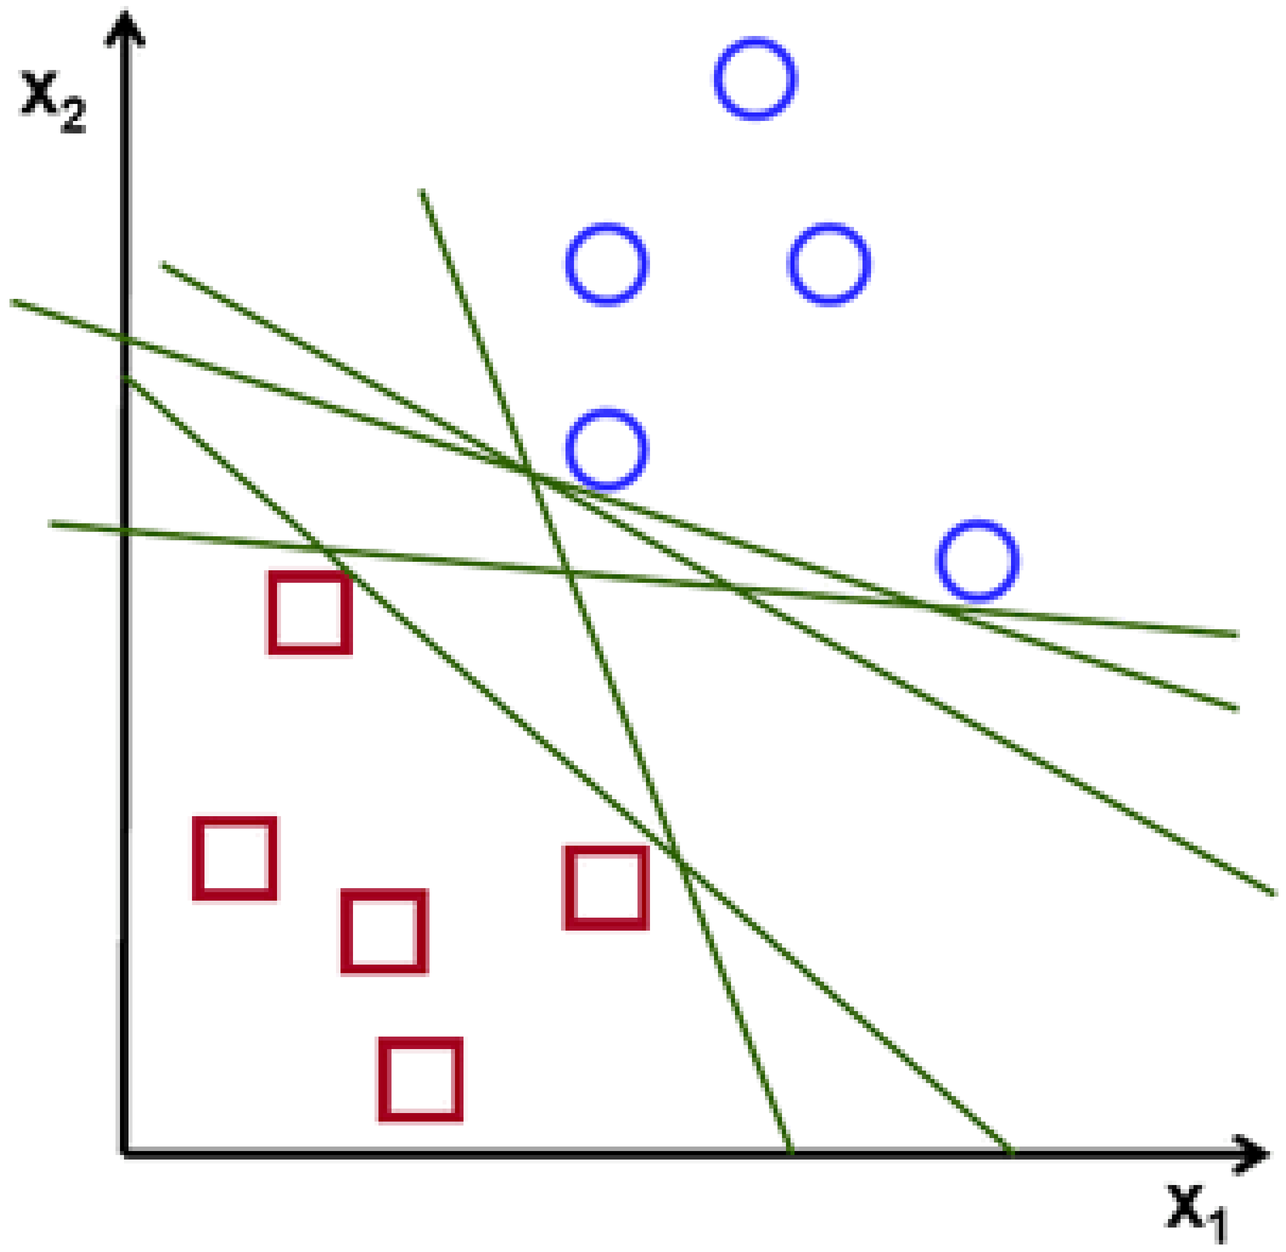

SUPPORT VECTOR MACHINE (SVM)

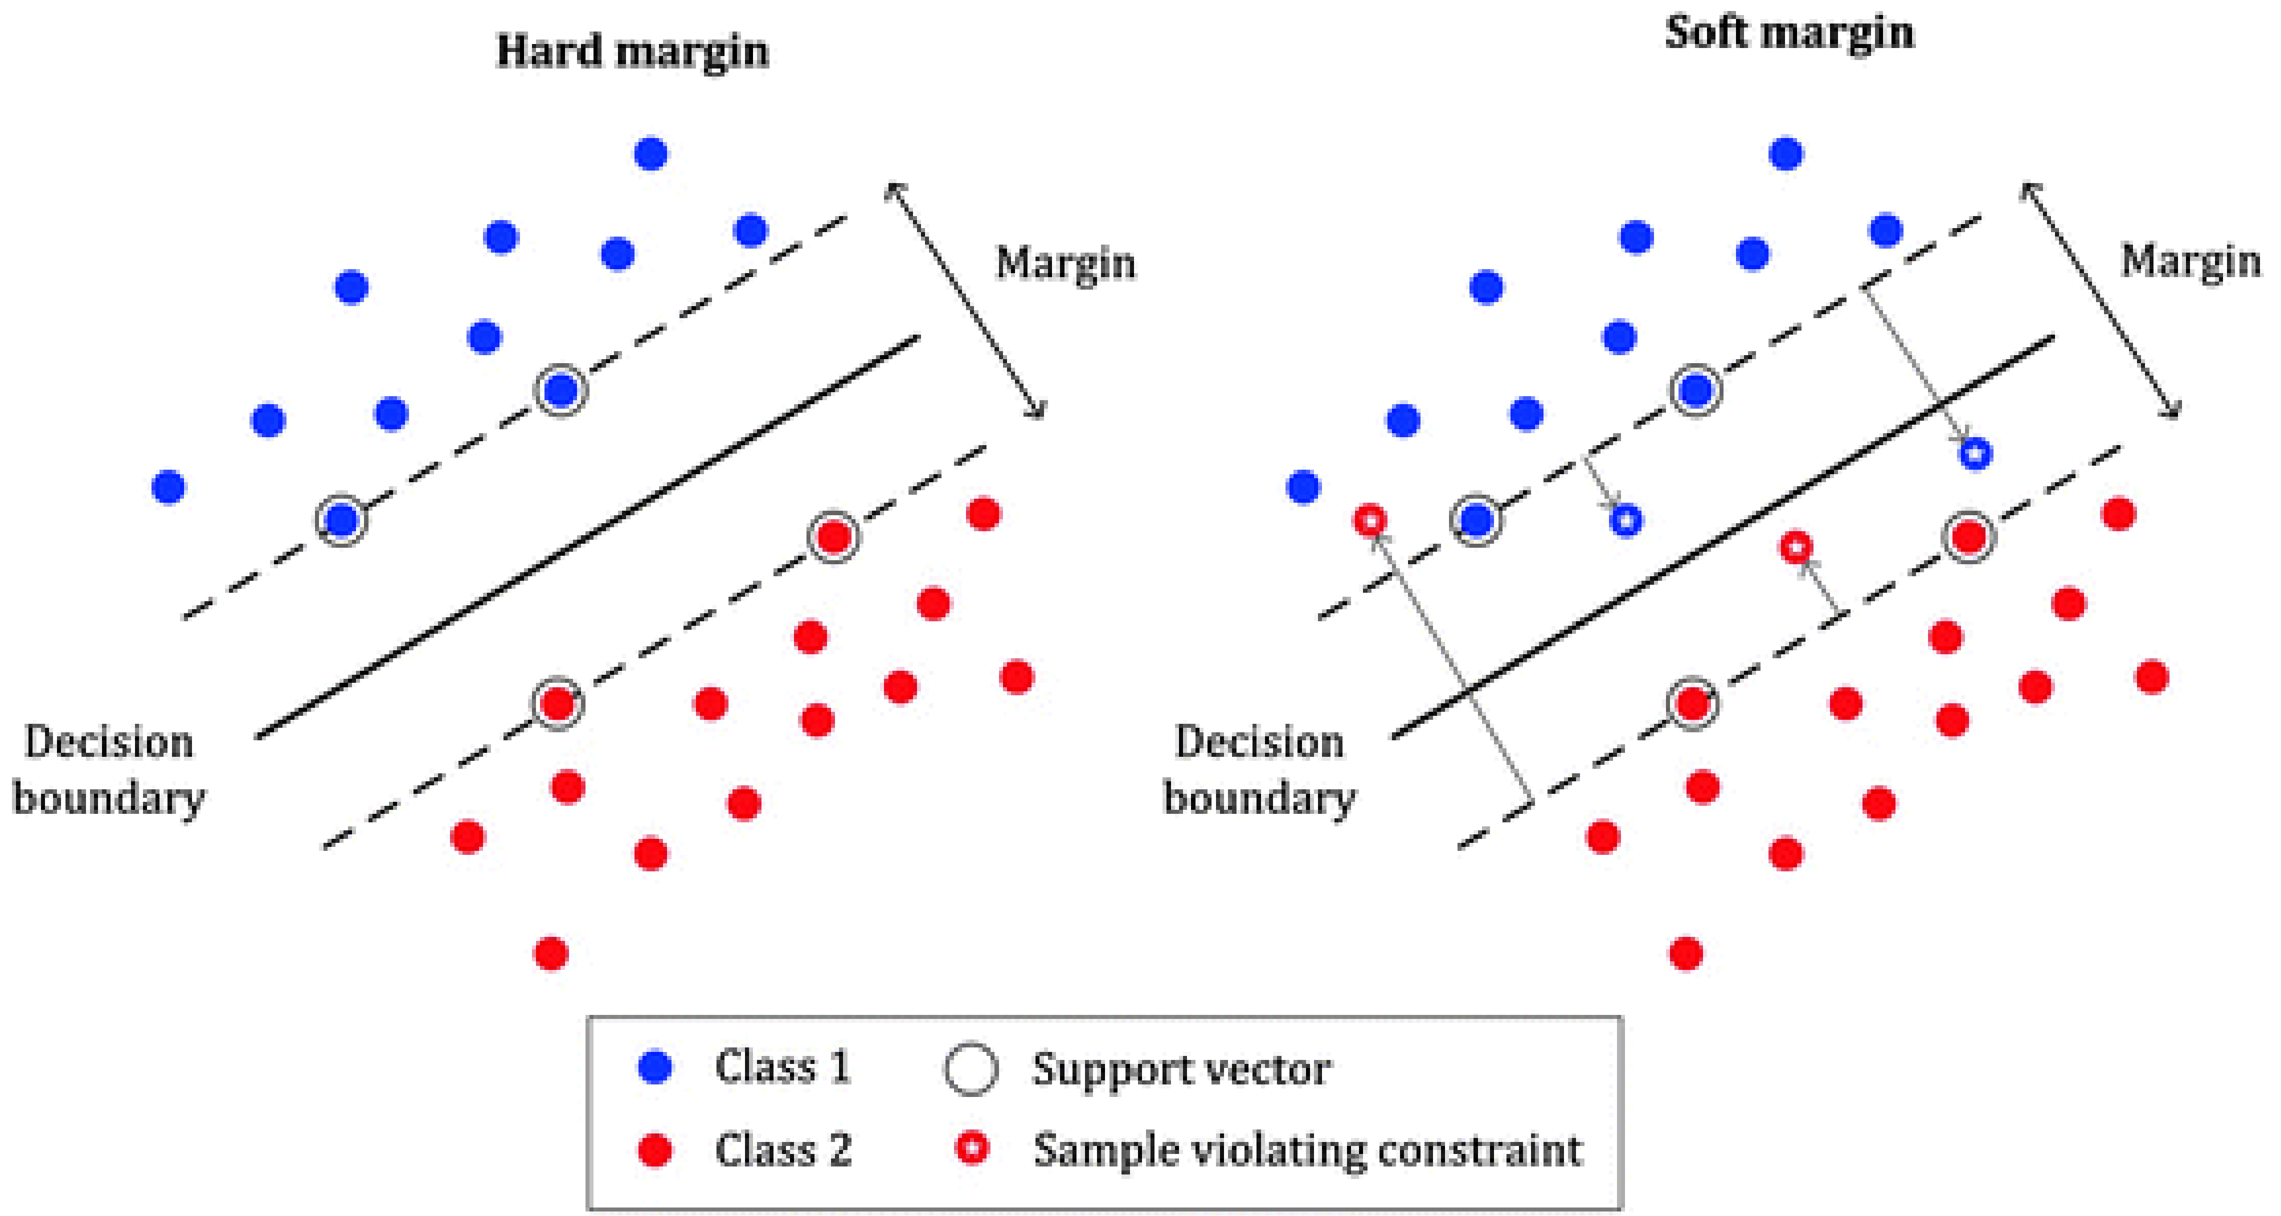

Our objective is to find a linear boundary that has the maximum margin, i.e the maximum distance between data points of both classes. Maximizing the margin distance provides some reinforcement so that future data points can be classified with more confidence.

We can force the SVM to maximize the distance to any point in any of the sets or add a loss function to allow but penalize some violations of the boundary

SUPPORT VECTOR MACHINE (SVM)

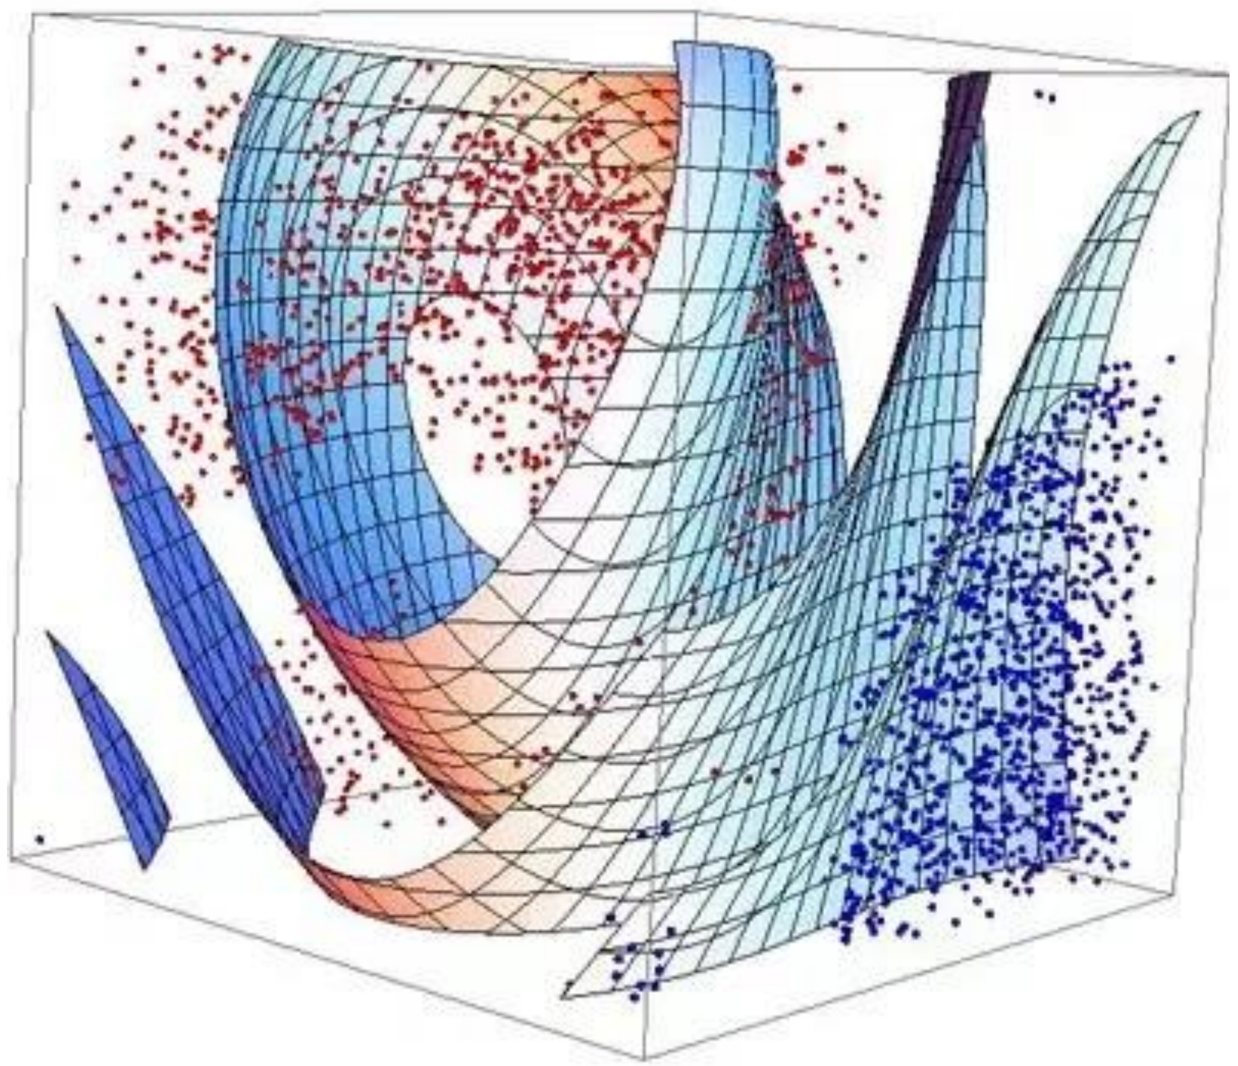

What if there is no such boundary?

We can usually transform the underlying datapoints in such a way that a linear boundary exists. Then we revert that transformation and our boundary stops being linear, but is still effective for classification.

This is called a Kernel trick and usually requires very good understanding of the data

SVM - SUMMARY

+

Can Model NonLinear Boundaries (if nonlinear version is used)

Unlikely to overfit

−

Memory intensive: is more successful on small dataset